



SECURE
ENGAGING
EVERYWHERE



Cumulus Meetup 2018.10.10 - Zurich

Network automation at Kudelski – Romain Aviolat

Agenda

Romain Aviolat

Cloud Infrastructure Expert

- Swiss product
- Mountains
- Reading, Cinema, Music
- Open Source advocate
- Void warranties
- Hardware hacking



Kudelski Group

- +60 years
- +3K employees on 5 continents
- 200M+ annual R&D investment
- DigitalTV (Content protection)
- Public Access
- Cyber Security
- IoT



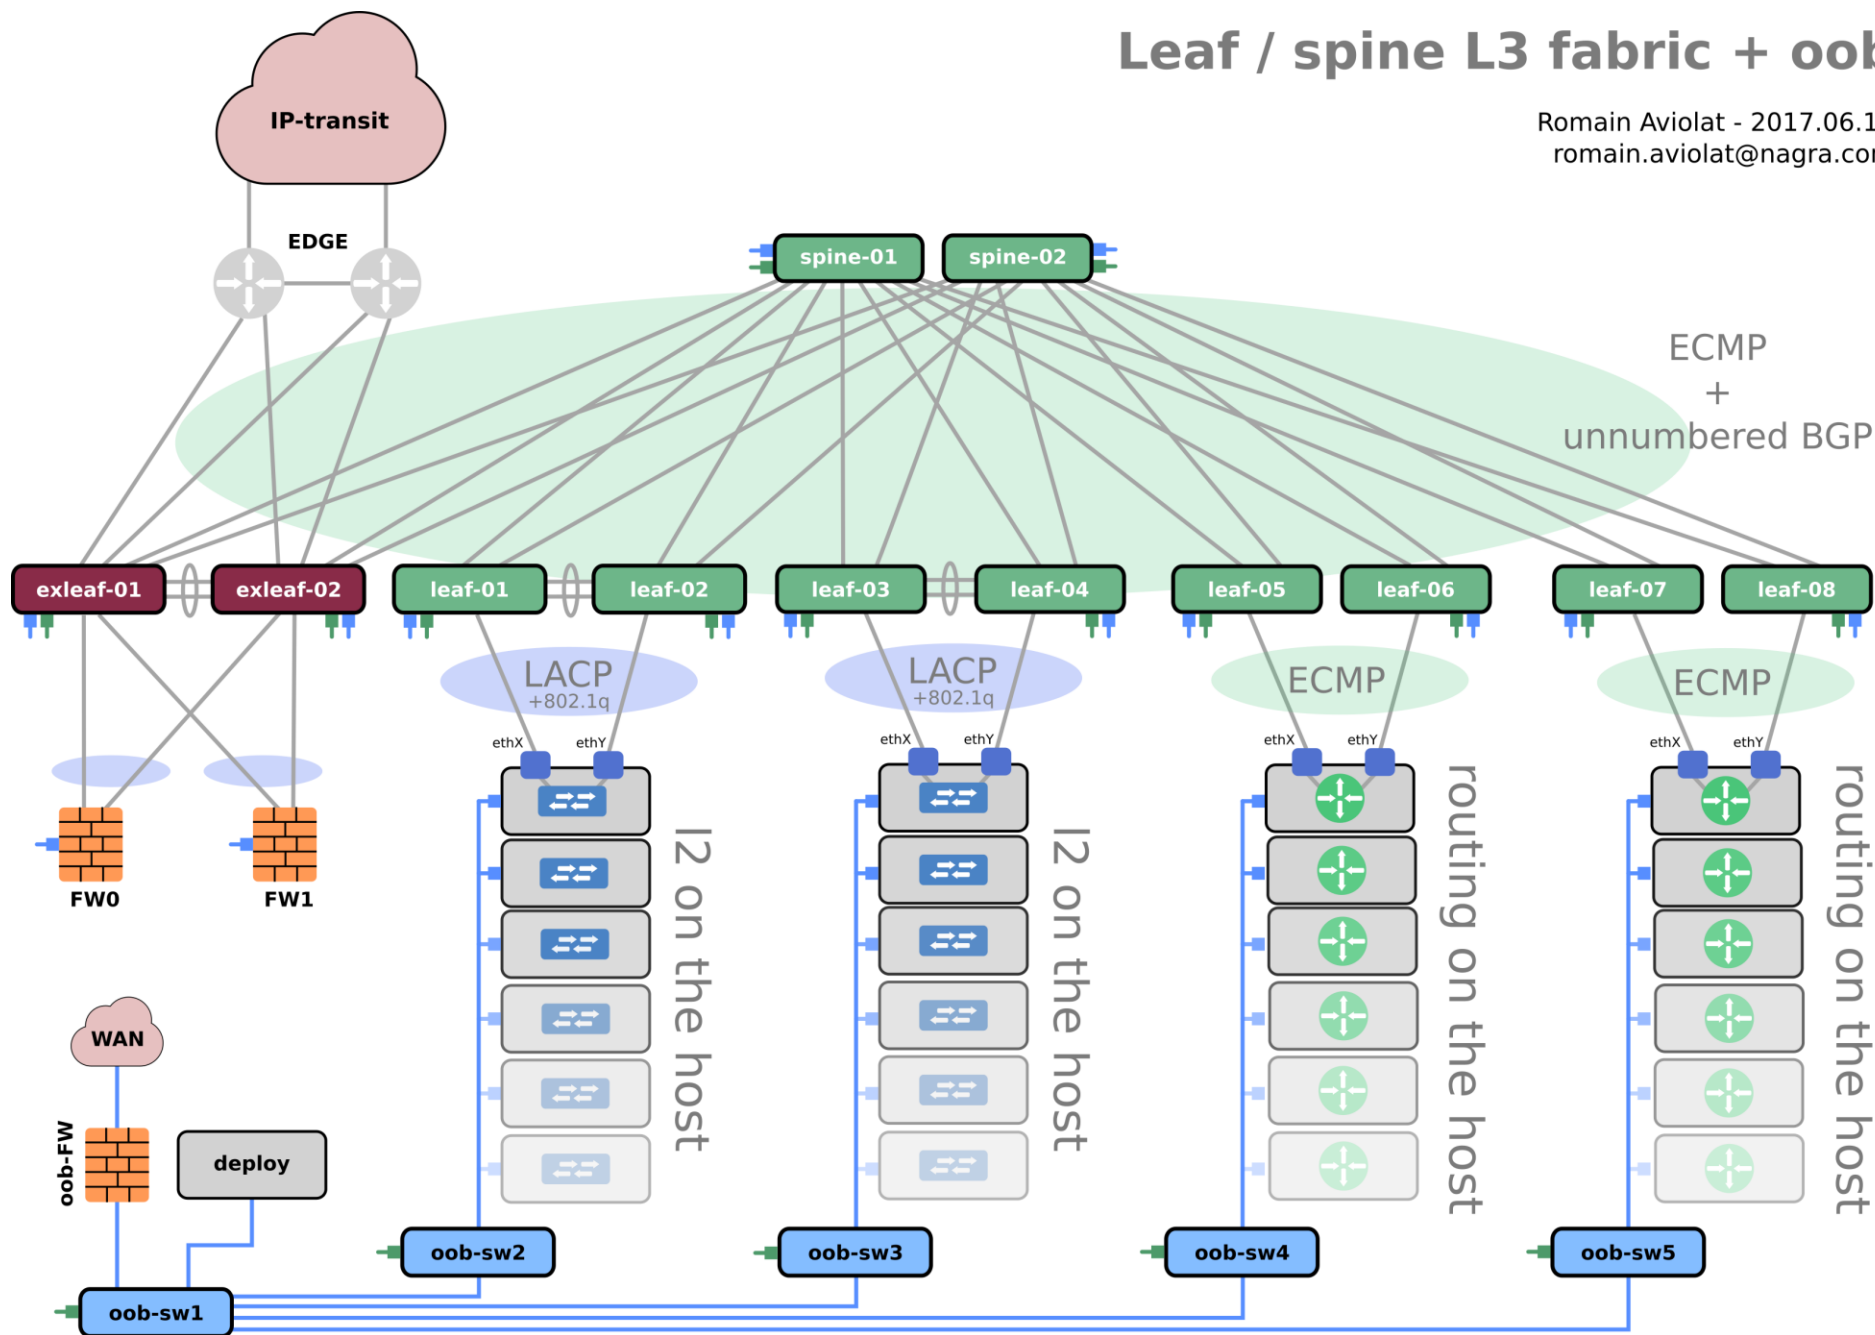
en.wikipedia.org/wiki/Kudelski_Group

Team / missions (IT cloud infrastructure)

- Design / Build / Operate the foundation of our new IT infrastructure, that serves the whole group.
- Must be flexible and aaS as possible for R&D teams
- Silo-less mode / full stack
- 7 engineers
- DevOps (Sec) principle

Techno stack

- IaaS: OpenStack
- Storage (+object): CEPH
- VMware cloud: NSX / vRA
- Baremetal deployment: MaaS (Canonical)
- Network: Whitebox + cumuluslinux
- Containers: K8s, docker
- IaaS: Ansible, Saltstack
- CI/CD + versioning: GitLab
- Monitoring: Prometheus, Grafana, Kibana



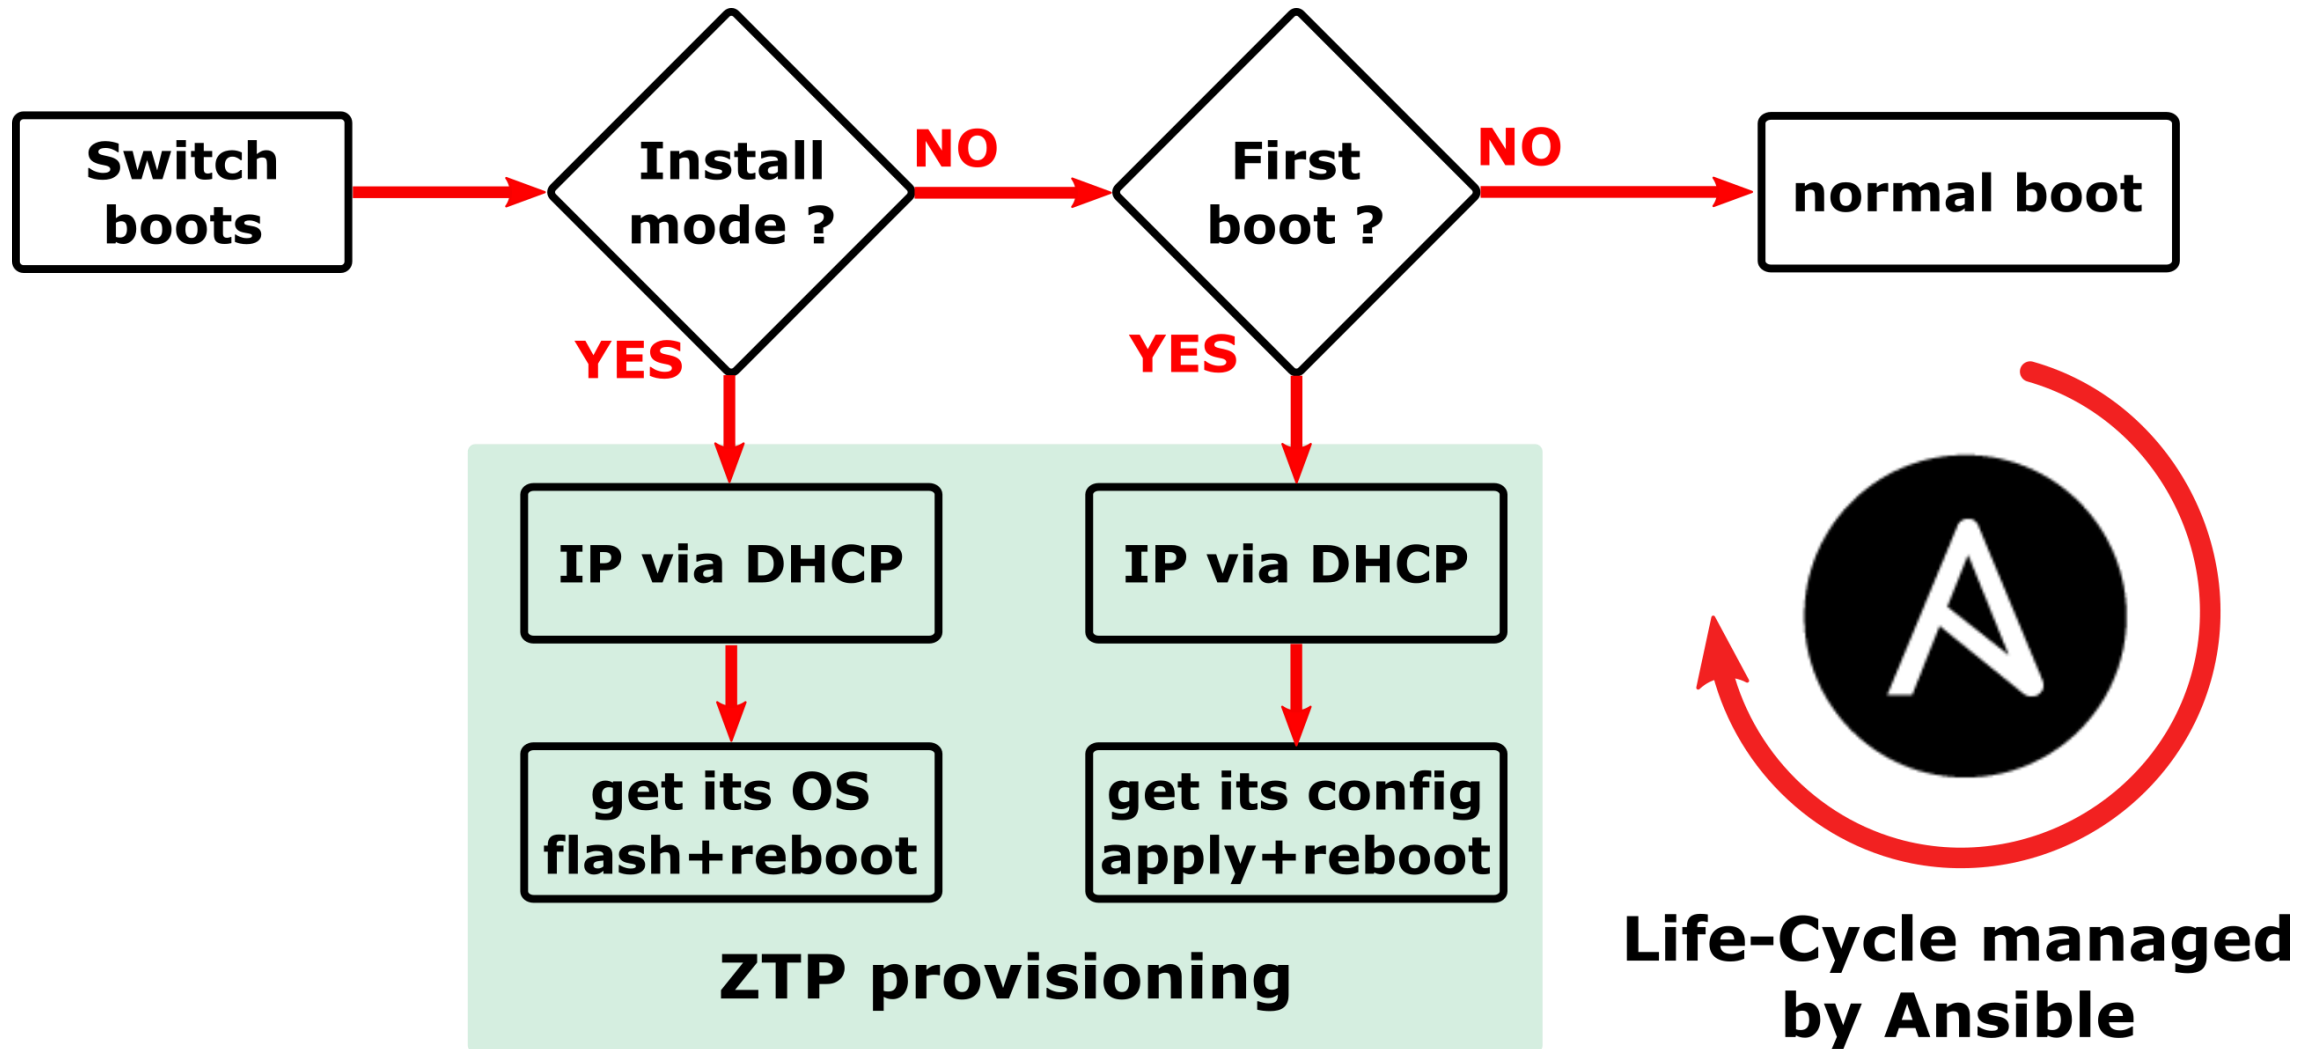
Network features

- Layer 2 virtualization using eVPN
- L3 down to the host (Bare-metal infrastructure)
 - 1 IP (v4/v6) per host and per service
 - Anycast v4 + v6 for services HA
- VRFs

Do it as-code or don't do it

- We automate everything (as in everything as possible), using:
 - APIs (custom tools if doesn't exist)
 - Config-mgmt tools (Ansible, SaltStack)
- It's a strong criteria when selecting hardware or software
- Starting to do that also to lower operational tasks
- We of course also apply this principle for network appliances

Devices life-cycle



Ansible variables against templates

```
interfaces:
  swp1:
    alias: storage-dc1r02n01
    vlans:
      - ipmi
    pvid: 2048
  swp2:
    alias: storage-dc1r02n02
    vlans:
      - ipmi
    pvid: 2048
  swp3:
    alias: storage-dc1r02n03
    vlans:
      - ipmi
    pvid: 2048
  swp4:
    alias: storage-dc1r02n04
    vlans:
      - ipmi
    pvid: 2048
```

```
{% if interfaces is defined %}
{% for port, value in interfaces.items() %}
auto {{ port }}
iface {{ port }} {% if value and 'address'

{% if value and 'mtu' in value %}
    mtu {{ value['mtu'] }}
{% endif %}
{% if value and 'link-speed' in value %}
    link-speed {{ value['link-speed'] }}
{% endif %}
```

```
auto swp2
iface swp2
    bridge-vids 1536
    bridge-pvid 2048
    alias storage-dc1r02n02

auto swp3
iface swp3
    bridge-vids 1536
    bridge-pvid 2048
    alias storage-dc1r02n03

auto swp1
iface swp1
    bridge-vids 1536
    bridge-pvid 2048
    alias storage-dc1r02n01
```

- Routing engine, interfaces, ...
- Use variable groups to maintain consistency
- We don't use custom modules

Git for versioning

```
swp5:
  link-autoneg: true
  link-speed: 1000
+ swp8:
+   alias: ipanema-lan
+   access: 1538
+   link-speed: 1000
+   link-autoneg: false
swp35:
  alias: service-dc1r01n01
  l3host: true
```

```
commit bc889ac62847d03f38ad6db26edc47cfd5e584fe (origin/
Author: Aviolat Romain <romain.aviolat@nagra.com>
Date: Thu Oct 4 08:37:18 2018 +0200

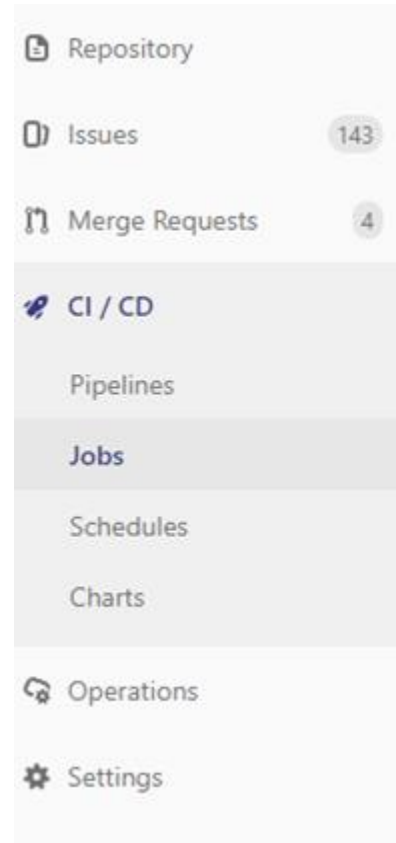
    fabric: add configs for Ipanema

commit 5b2731598a40329e4063cd7a9c5a2d3c09818821
Author: Aviolat Romain <romain.aviolat@nagra.com>
Date: Thu Oct 4 08:36:40 2018 +0200

    doc: vlan: update to reflect last changes
```

GitLab / Github / CI/CD platform

- Don't run your code locally on your machine
- Describe your work inside issues
- Ansible code / Pipelines are triggered by commits
- Merge / Pull requests to push in production
- History of all Ansible runs



```
TASK [dhcp_server : define static dhcp mappings] *****
ok: [adminsw-dc1r01n01. ]
Wednesday 10 October 2018 13:45:06 +0200 (0:00:03.539) 0:03:58.566 *****

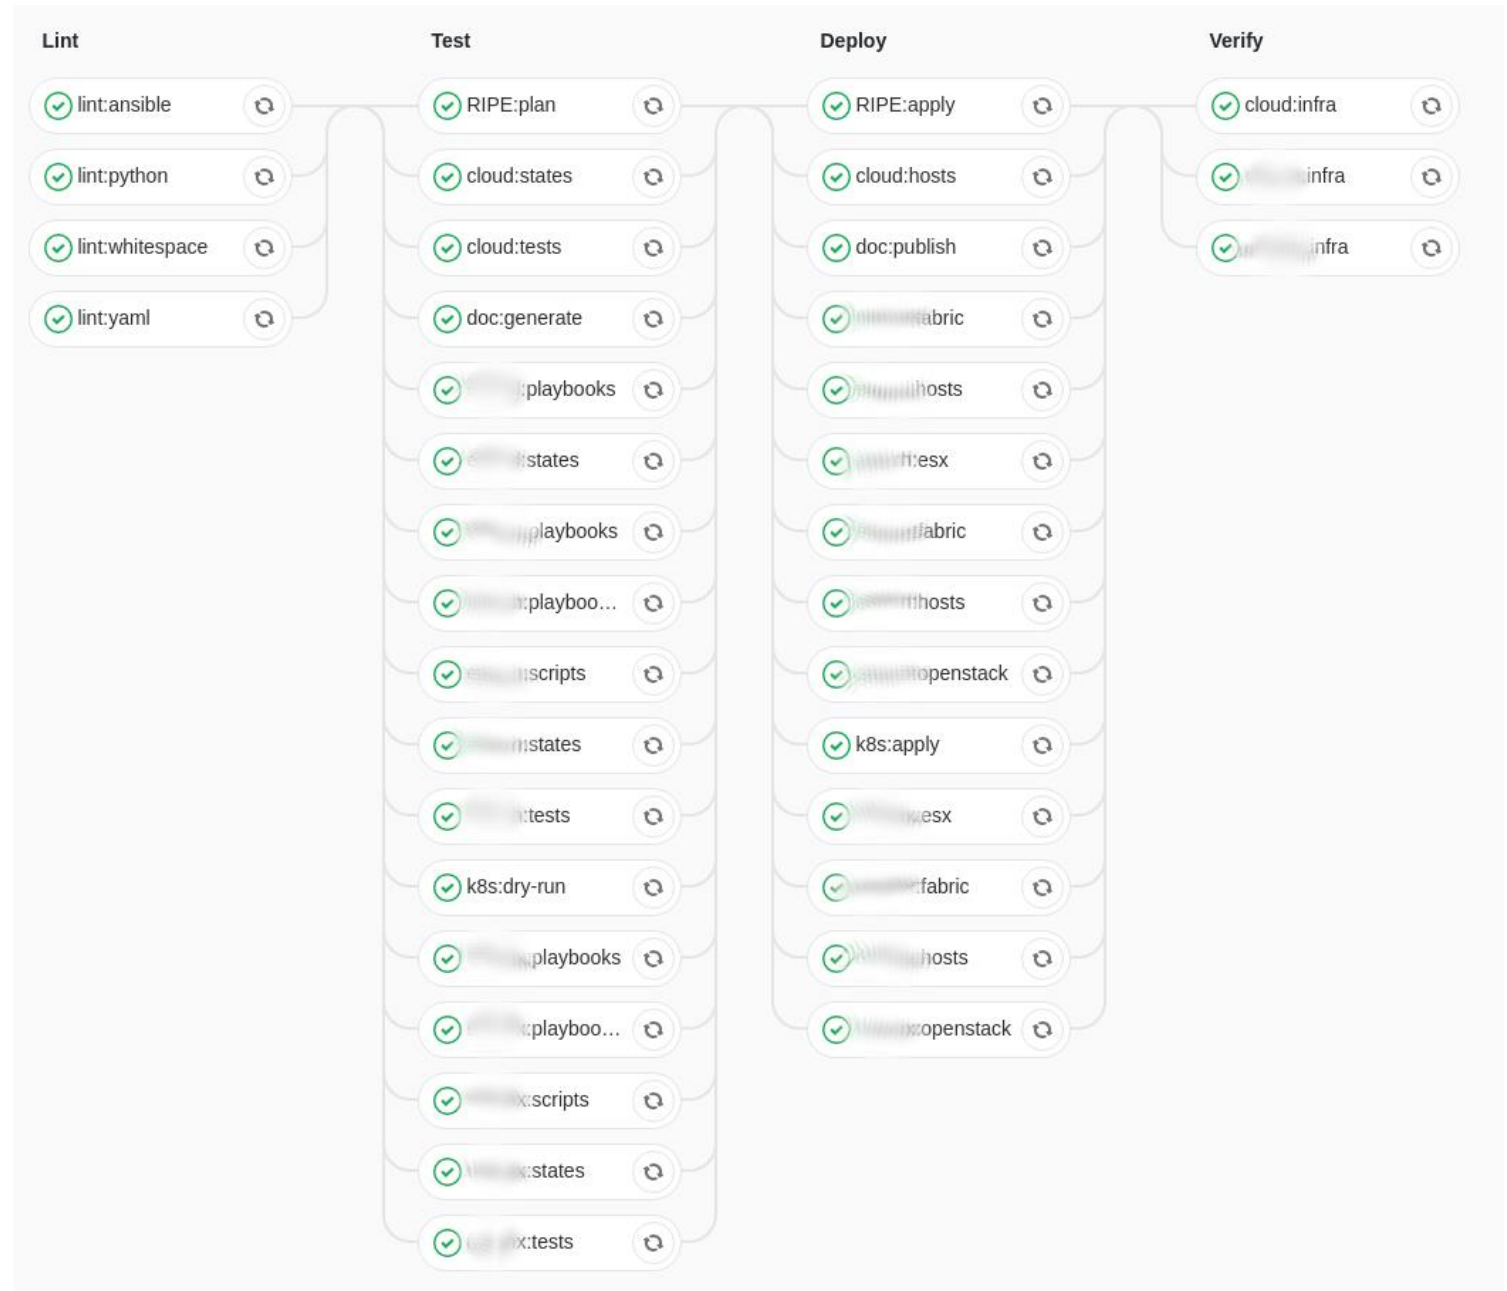
TASK [dhcp_server : set dhcp options] *****
ok: [adminsw-dc1r01n01. ]
Wednesday 10 October 2018 13:45:09 +0200 (0:00:03.399) 0:04:01.966 *****

TASK [dhcp_server : configure dnsmasq] *****
ok: [adminsw-dc1r01n01. ]

PLAY RECAP *****
adminsw-dc1r01n01. : ok=33 changed=0 unreachable=0 failed=0
adminsw-dc1r02n01. : ok=26 changed=0 unreachable=0 failed=0
adminsw-dc1r03n01. : ok=26 changed=0 unreachable=0 failed=0
exleaf-dc1r01n01. : ok=32 changed=1 unreachable=0 failed=0
exleaf-dc1r01n02. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r02n01. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r02n02. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r03n01. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r03n02. : ok=32 changed=1 unreachable=0 failed=0
spine-dc1r01n01. : ok=32 changed=1 unreachable=0 failed=0
spine-dc1r01n02. : ok=32 changed=1 unreachable=0 failed=0
```


GitLab CI/CD

- Multiple stages:
 - Lint (very important)
 - Test
 - Deploy
 - Verify
- Multiple Job:
 - Fabric
 - DNSes
 - Hosts
 - ...




Follow sw-dev. best-practices

- Guys been doing that for years (at scale)
 - Unit tests
- Multiple environments
 - Dev, Staging, Production
- Code review / Four-eyes review
 - Don't push in production what you coded
 - Ask someone to do it (he'll become responsible)
- CI / CD pipeline



Numbers

- We merged 2500 time in production since August 2016



- Managing 3 DataCenters this way
- 55 network devices, 100's servers
- Repetability is key

Automating operational tasks

- Reducing operational load
- Decrease human errors
 - 20 appliances to upgrade it's very likely that you'll do mistakes
- Decrease number of suicides
- Appliances auto-upgrade

Upgrade my DataCenter ! (scary part)

- Custom set of playbooks / scripts developed internally to upgrade au whole DC
- APT upgrades only for now (no binary upgrades)

```
- hosts: switches
  user: cumulus
  serial: 1

  tasks:

    #####
    #### check if the switch needs to be upgraded or not

    - name: Register the OS version
      shell: grep "VERSION_ID=" /etc/os-release | cut -d "=" -f 2
      register: current_version
      tags: register

    - name: Check if switch needs to be upgraded or not
      block:
        - debug:
            msg: "Switch needs to be upgraded from {{ current_version.stdout }}"


        #####
        #### pre-flight checks applied to all the fabric switches

        # Execute the consistency checks on all the switches
        - include_tasks: fabric_consistency.yml

        - pause:
            prompt: "Make sure that the peer has become master and hit enter"
```


Time-Series-Based Monitoring

- Collect data then create relevant dashboards alerts
 - Make them evolve over time, data will be here anyway
- Powerful queries language, combine metrics
- Prometheus / Grafana
- Alerting
- SNMP



Staging / LABs

- Complicated to achieve with physical hardware
- With virtualization it's now “easy” to simulate a complete network environment
- KVM, Virtualbox / Vagrant
- Some vendors directly provide a VM for their OS



Challenges

- Optimize the Ansible code to make it fast (but not too fast...)
- i40e Linux drivers
 - broken on Ubuntu 14.04, had to hack it a bit
 - Post-spectre memory leak
- Non x86/amd64 platform (oob)
 - Binaries (in)compatibility