



SECURE
ENGAGING
EVERYWHERE



Cumulus Meetup 2019.03.26 - Lausanne


Network automation at Kudelski – Romain Aviolat

Agenda

- Network automation
- DevOps
- Return of experience (>3y)

About ME

- (social) Nerd
 - Open-source advocate
 - Linux / Network / Sec
 - Hardware hacker (μ C, Drones, toasters, ...)
 - Endurance-sports enthusiast
 - Don't like products
- Team lead infrastructure team (IT structure)
 - Small cross-functional team
- Joined the group 5 years ago



Kudelski Group

- +60 years
- +3K employees on 5 continents
- 200M+ annual R&D investment
- DigitalTV (Content protection)
- Public Access
- Cyber Security
- IoT



en.wikipedia.org/wiki/Kudelski_Group

The team

(thanks guys)

“Provide multi-tenant and as-a-service infrastructure for the whole **K** group, based on new foundations. Inside Tier3+ Datacentres, close to our users (/clients)”

- Infrastructure (IT)
- DataCenters
 - hosting / collocation
 - IP-transit (LIR)
- Private-Cloud
 - OpenStack + Vmware
 - Baremetal / DataScience
- Services
 - DevOps tooling

- + Cross functional
- + Everybody's hands-on ! / no golden-fingers
- + Generalize as-code for everything
- + Commit often, perfect later, push once





Team mojo – “Do it as-code or don’t do it”

- We automate everything (as in everything as possible), using:
 - Config-mgmt tools (Ansible, SaltStack)
 - APIs (custom tools if doesn’t exist)
 - DNSes, RIPE object, IaaS, ...
- It’s a strong criteria when selecting hardware or software
- Starting to do that also to lower operational tasks
- We of course also apply this principle for network appliances

Refresher: Traditional model vs DevOps

OLD WAY:



- + Follow software dev best practices
- + Try to apply the same practices for infra
- + Never ending story



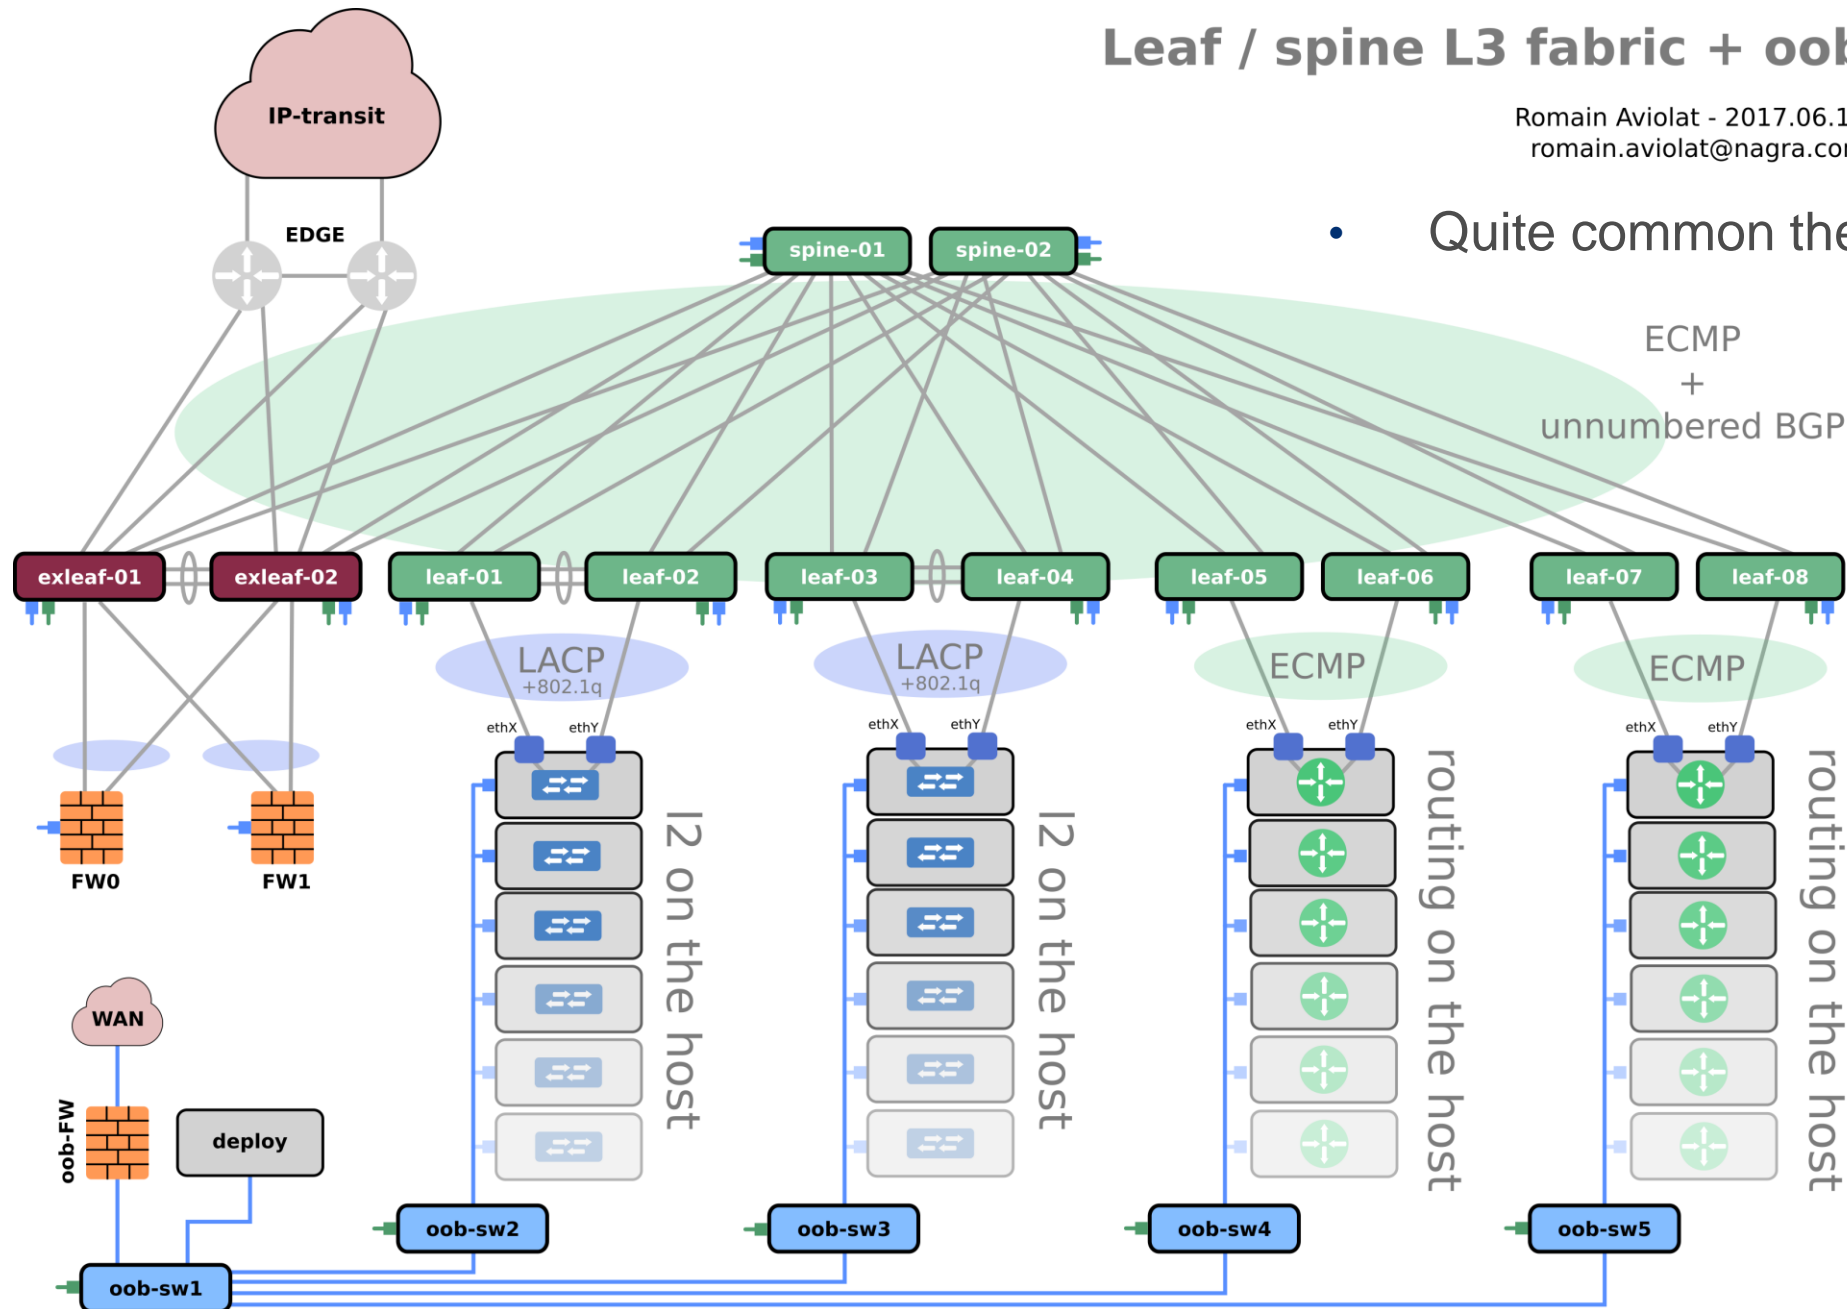
Techno stack

- IaaS: OpenStack
- Storage (+object): CEPH
- VMware cloud: NSX / vRA
- Baremetal deployment: MaaS (Canonical)
- Network: Whitebox + cumuluslinux
- Containers: K8s, docker
- IaaS: Ansible, Saltstack
- CI/CD + versioning: GitLab
- Monitoring: Prometheus, Grafana, Kibana, Grafana

Leaf / spine L3 fabric + oob

Romain Aviolat - 2017.06.19
romain.aviolat@nagra.com

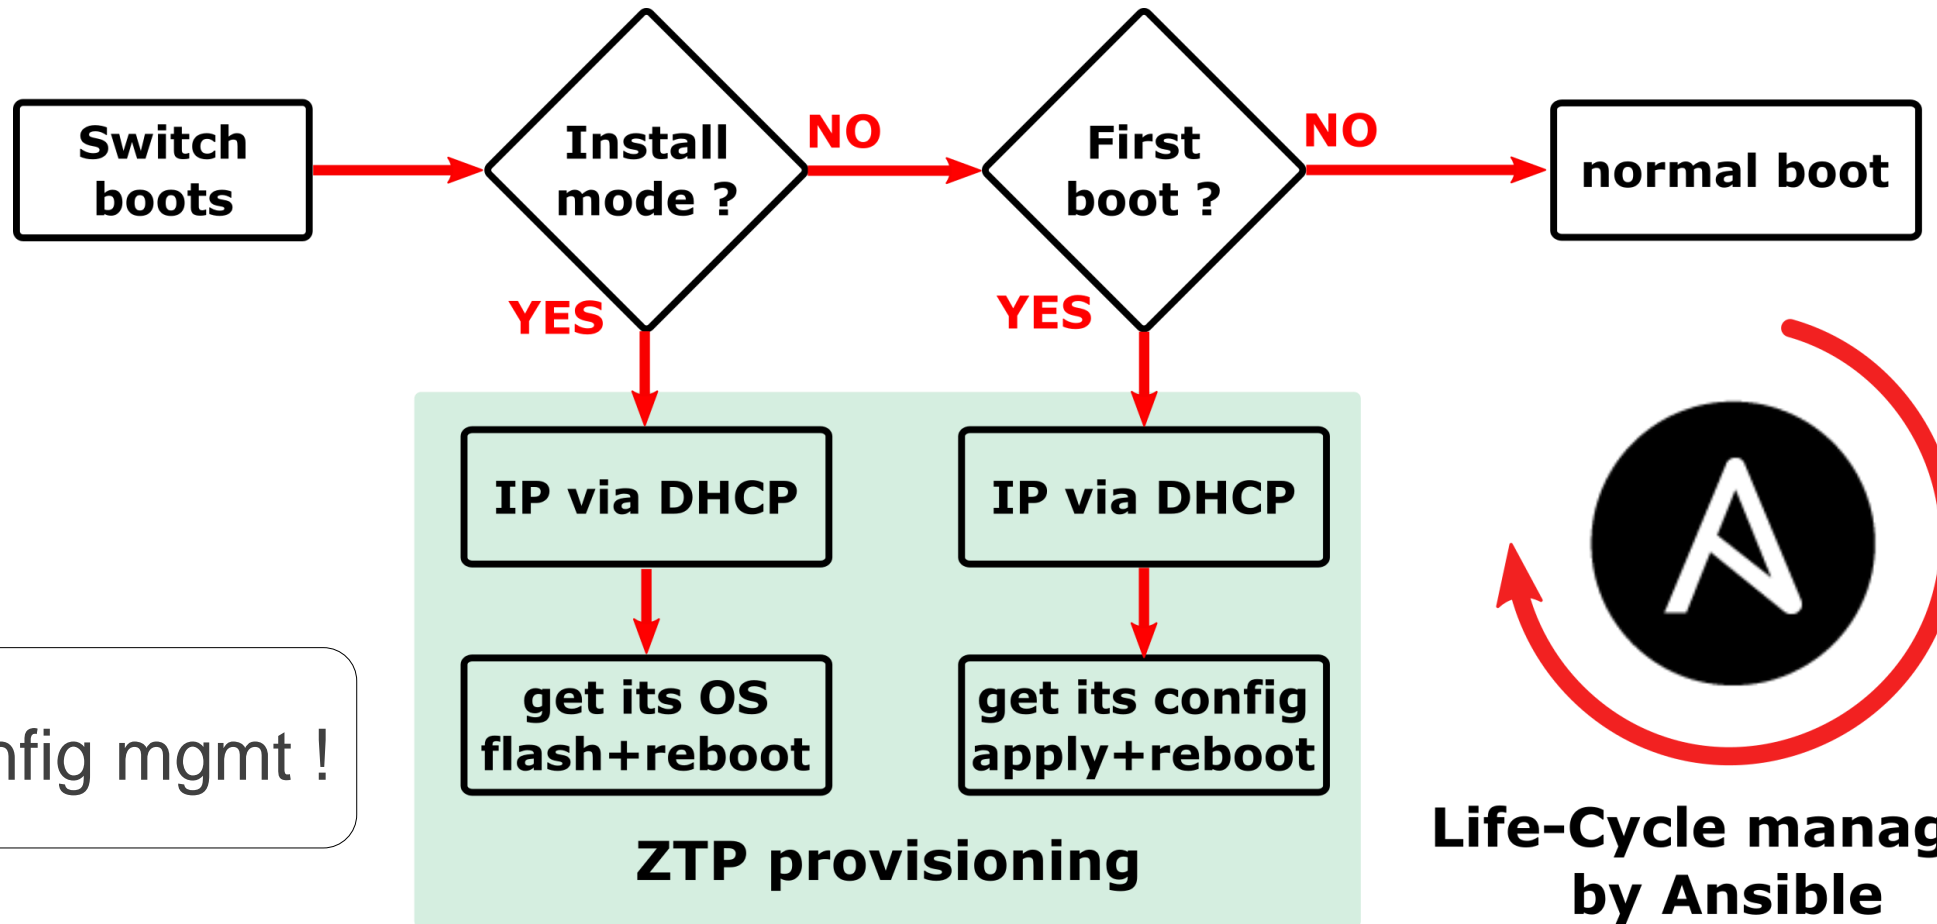
- Quite common these days




Network features

- 1 / 10 / 40GbE
- Layer 2 virtualization using eVPN
- L3 down to the host (Bare-metal infrastructure)
 - 1 IP (v4/v6) per host and per service
 - Anycast v4 + v6 for services HA
 - FRR
- VRFs for data / control plane

Devices life-cycle



- ZTP != config mgmt !



Config management using Ansible

```
interfaces:
  swp1:
    alias: storage-dc1r02n01
    vlans:
      - ipmi
    pvid: 2048
  swp2:
    alias: storage-dc1r02n02
    vlans:
      - ipmi
    pvid: 2048
  swp3:
    alias: storage-dc1r02n03
    vlans:
      - ipmi
    pvid: 2048
  swp4:
    alias: storage-dc1r02n04
    vlans:
      - ipmi
    pvid: 2048
```

```
{% if interfaces is defined %}
{% for port, value in interfaces.items() %}
auto {{ port }}
iface {{ port }} {% if value and 'address' %}
    address {{ value['address'] }}
{% endif %}
{% if value and 'mtu' in value %}
    mtu {{ value['mtu'] }}
{% endif %}
{% if value and 'link-speed' in value %}
    link-speed {{ value['link-speed'] }}
{% endif %}
{% endfor %}
```

```
auto swp2
iface swp2
    bridge-vids 1536
    bridge-pvid 2048
    alias storage-dc1r02n02

auto swp3
iface swp3
    bridge-vids 1536
    bridge-pvid 2048
    alias storage-dc1r02n03

auto swp1
iface swp1
    bridge-vids 1536
    bridge-pvid 2048
    alias storage-dc1r02n01
```

- Routing engine, interfaces, ...
- Use variable groups to maintain consistency
- We don't use custom modules
- Repeatability is keys (really)

Git for versioning

- Collaboration
- Track changes

```
swp5:
    link-autoneg: true
    link-speed: 1000
+ swp8:
+   alias: ipanema-lan
+   access: 1538
+   link-speed: 1000
+   link-autoneg: false
swp35:
    alias: service-dc1r01n01
    l3host: true
```

```
commit bc889ac62847d03f38ad6db26edc47cfd5e584fe (origin/
Author: Aviolat Romain <romain.aviolat@nagra.com>
Date: Thu Oct 4 08:37:18 2018 +0200
```

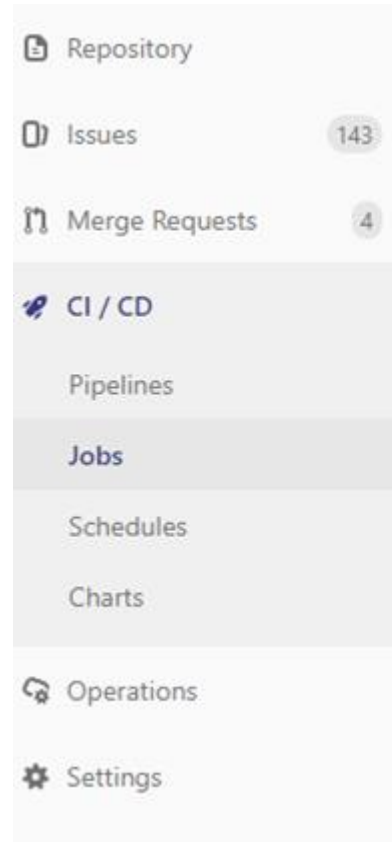
```
    fabric: add configs for Ipanema
```

```
commit 5b2731598a40329e4063cd7a9c5a2d3c09818821
Author: Aviolat Romain <romain.aviolat@nagra.com>
Date: Thu Oct 4 08:36:40 2018 +0200
```

```
    doc: vlan: update to reflect last changes
```

GitLab / Github / CI/CD platform

- Don't run your code locally on your machine
- Describe your work inside issues
- Ansible code / Pipelines are triggered by commits
- Merge / Pull requests to push in production
- History of all Ansible runs



```
TASK [dhcp_server : define static dhcp mappings] *****
ok: [adminsw-dc1r01n01. ]
Wednesday 10 October 2018 13:45:06 +0200 (0:00:03.539) 0:03:58.566 *****

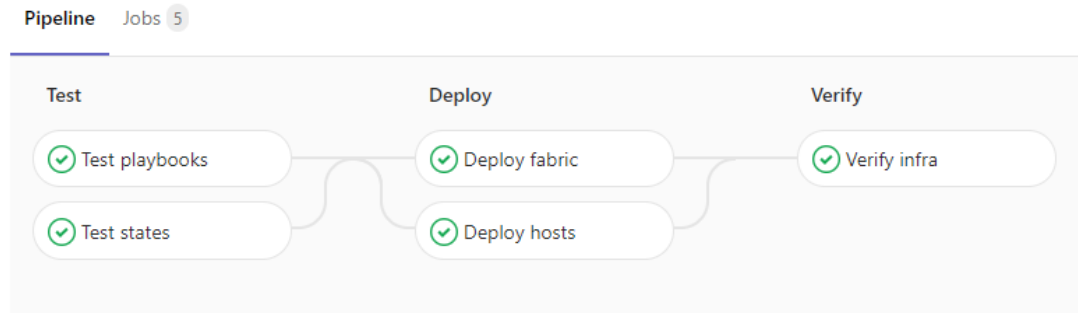
TASK [dhcp_server : set dhcp options] *****
ok: [adminsw-dc1r01n01. ]
Wednesday 10 October 2018 13:45:09 +0200 (0:00:03.399) 0:04:01.966 *****

TASK [dhcp_server : configure dnsmasq] *****
ok: [adminsw-dc1r01n01. ]

PLAY RECAP *****
adminsw-dc1r01n01. : ok=33 changed=0 unreachable=0 failed=0
adminsw-dc1r02n01. : ok=26 changed=0 unreachable=0 failed=0
adminsw-dc1r03n01. : ok=26 changed=0 unreachable=0 failed=0
exleaf-dc1r01n01. : ok=32 changed=1 unreachable=0 failed=0
exleaf-dc1r01n02. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r02n01. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r02n02. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r03n01. : ok=32 changed=1 unreachable=0 failed=0
leaf-dc1r03n02. : ok=32 changed=1 unreachable=0 failed=0
spine-dc1r01n01. : ok=32 changed=1 unreachable=0 failed=0
spine-dc1r01n02. : ok=32 changed=1 unreachable=0 failed=0
```

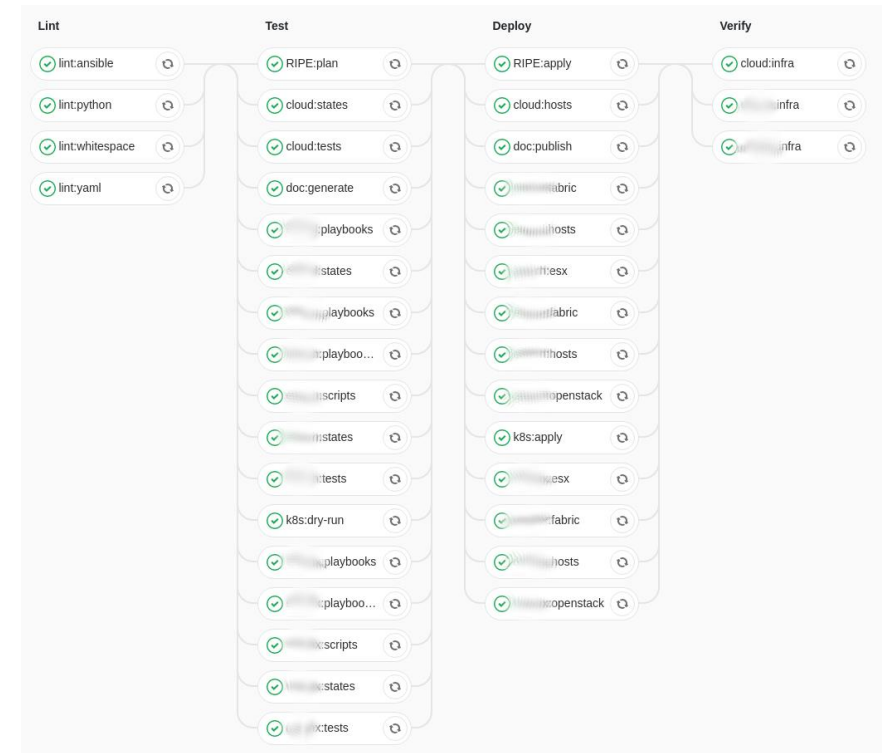

Start simple !

2016



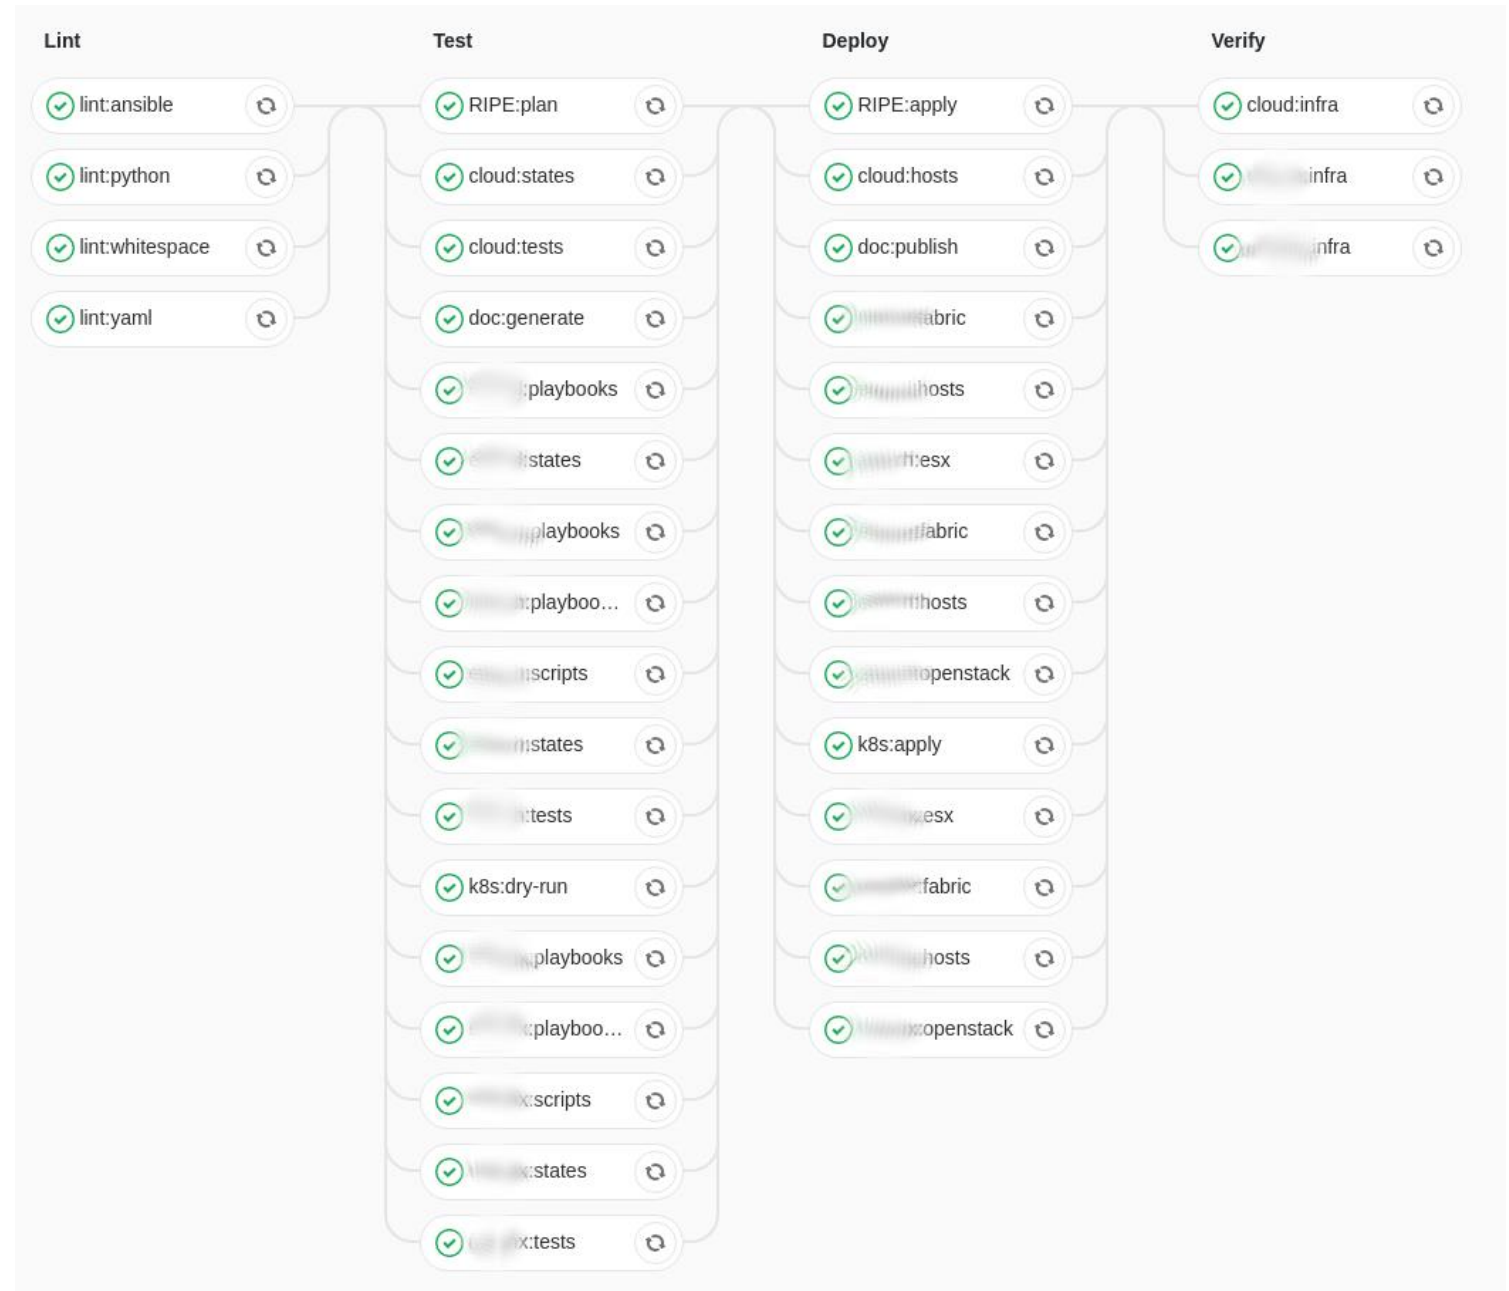
- Start with low-hanging fruits first

2019




GitLab CI/CD

- Multiple stages:
 - Lint (very important)
 - Test
 - Deploy
 - Verify
- Multiple Job:
 - Fabric
 - DNSes
 - Hosts
 - ...

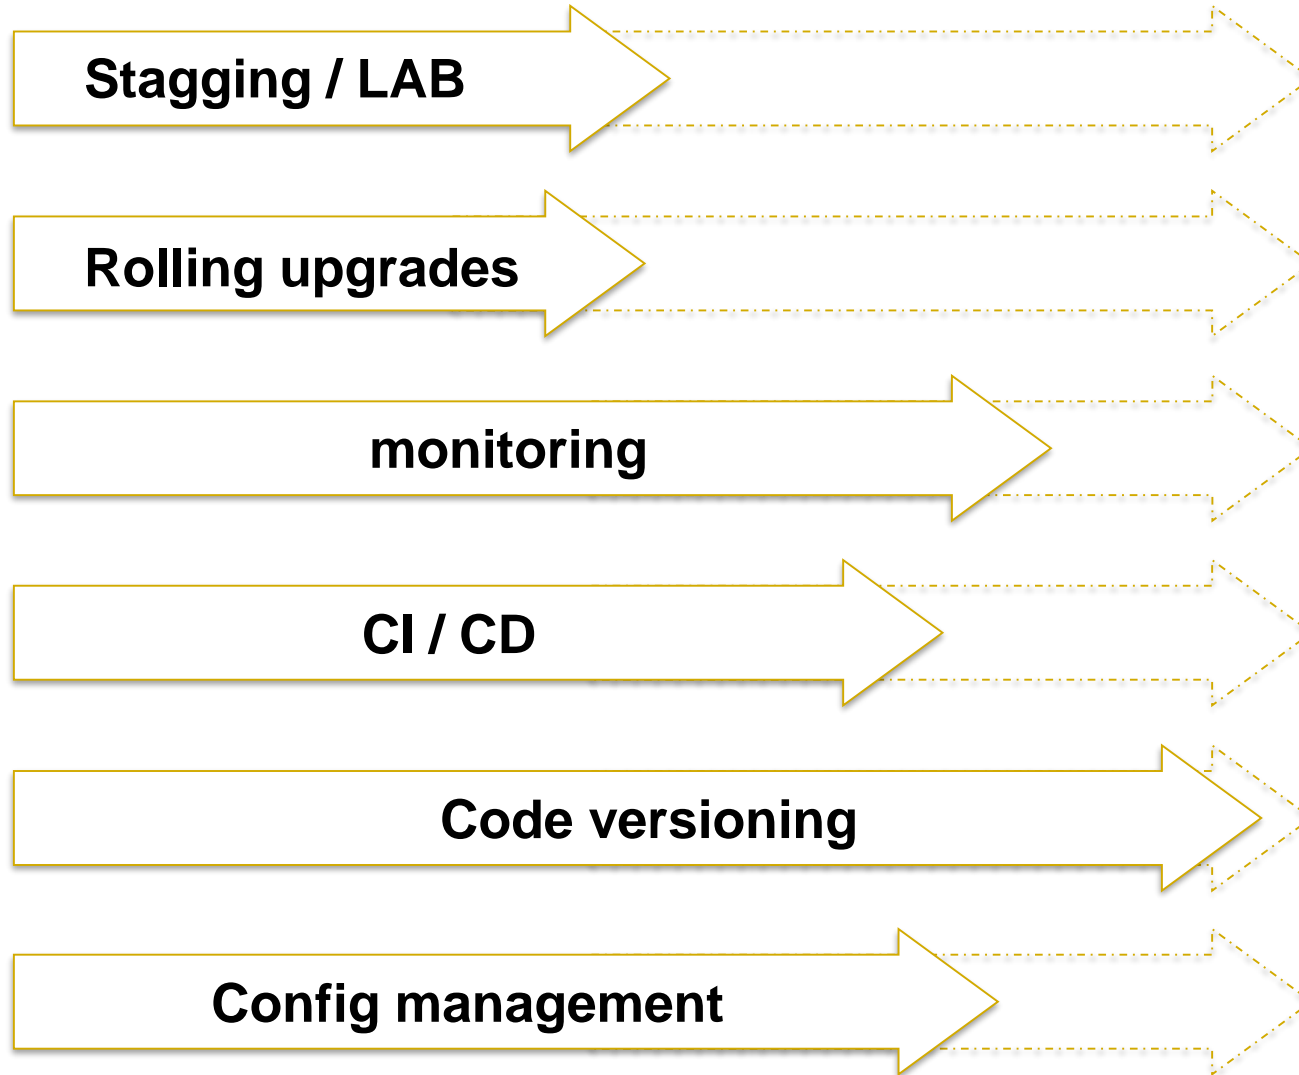


Follow sw-dev. best-practices

- Guys been doing that for years (at scale)
 - Unit tests
- Multiple environments (hard)
 - Dev, Staging, Production
- Code review / Four-eyes review
 - Don't push in production what you coded
 - Ask someone to do it (he'll become responsible)
- CI / CD pipeline



DevOpsness → never ending story




- + Document technical debt
- + Be honest (backlog)
- + Issue tracking
- + Direct communication



Numbers

- We merged 3000 time in production since August 2016

Pipelines for last year



- Managing 3 DataCenters this way
- ~60 network appliances

Automating operational tasks

- Reducing operational load
- Decrease human errors
 - 20 appliances to upgrade it's very likely that you'll do mistakes
- Appliances auto-upgrade

Upgrade my DataCenter ! (scary part)

- Custom set of playbooks / scripts developed internally to upgrade a whole DC
- APT upgrades only for now (no binary upgrades)

```
- hosts: switches
  user: cumulus
  serial: 1

  tasks:

    #####
    #### check if the switch needs to be upgraded or not

    - name: Register the OS version
      shell: grep "VERSION_ID=" /etc/os-release | cut -d "=" -f 2
      register: current_version
      tags: register

    - name: Check if switch needs to be upgraded or not
      block:
        - debug:
            msg: "Switch needs to be upgraded from {{ current_version.stdout }}"


        #####
        #### pre-flight checks applied to all the fabric switches

        # Execute the consistency checks on all the switches
        - include_tasks: fabric_consistency.yml

        - pause:
            prompt: "Make sure that the peer has become master and hit enter"
```

Staging / LABs

- Complicated to achieve with physical hardware
- With virtualization it's now “easy” to simulate a complete network environment
- KVM, Virtualbox / Vagrant
- Some vendors directly provide a VM for their OS



As-code != ultimate perfection

Filter files

ansible/inventory/group_vars

exleafs.us-phx.yml +2 -2

ansible/inventory/group_vars/exleafs.us-phx.yml

+ 2 - 2

Edit View file @ d1bae9de

...	...	@@ -307,13 +307,13 @@ vxlan1615:
307	307	mtu: 9050
308	308	
309	309	vxlan1615:
310		- bridge_access: 1615
	310	+ bridge_access: 1615
311	311	bridge_learning: off
312	312	vxlan_id: 1615
313	313	mtu: 9050
314	314	
315	315	vxlan1616:
316		- bridge_access: 1616
	316	+ bridge_access: 1616
317	317	bridge_learning: off
318	318	vxlan_id: 1616
319	319	mtu: 9050

Challenges

- Optimize the Ansible code to make it fast (but not too fast...)
- i40e Linux drivers
 - broken on Ubuntu 14.04, had to hack it a bit
 - Post-spectre memory leak
- Non x86/amd64 platform (oob)
 - Binaries (in)compatibility